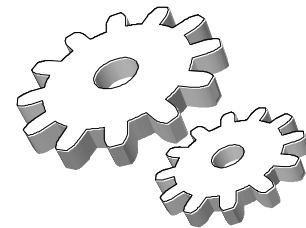


Simple Nested Loops Join

```
foreach tuple  $r \in R$  do  
    foreach tuple  $s \in S$  do  
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

Figure 12.4 Simple Nested Loops Join




Block Nested Loops Join


```
foreach block of  $B - 2$  pages of  $R$  do
  foreach page of  $S$  do {
    for all matching in-memory tuples  $r \in R\text{-block}$  and  $s \in S\text{-page}$ ,
    add  $\langle r, s \rangle$  to result
  }
```

Figure 12.5 Block Nested Loops Join

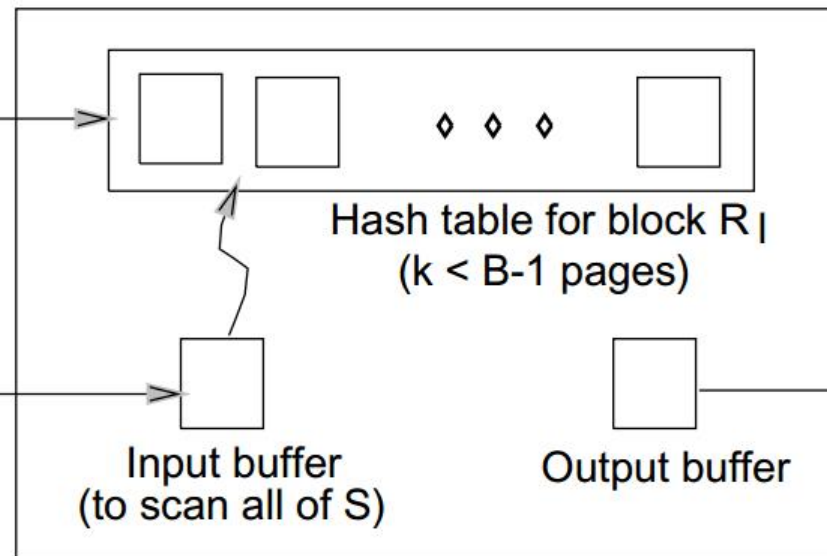
Buffer Usage in Block Nested Loops Join



Relations R and S




Disk




B main memory buffers

Join result



Disk


Figure 12.6 Buffer Usage in Block Nested Loops Join



Index Nested Loops Join

```
foreach tuple  $r \in R$  do  
    foreach tuple  $s \in S$  where  $r_i == s_j$   
        add  $\langle r, s \rangle$  to result
```

Figure 12.7 Index Nested Loops Join



```
proc smjoin( $R, S, 'R_i = S'_j$ )

if  $R$  not sorted on attribute  $i$ , sort it;
if  $S$  not sorted on attribute  $j$ , sort it;

 $Tr$  = first tuple in  $R$ ;                                // ranges over  $R$ 
 $Ts$  = first tuple in  $S$ ;                                // ranges over  $S$ 
 $Gs$  = first tuple in  $S$ ;                                // start of current  $S$ -partition


while  $Tr \neq eof$  and  $Gs \neq eof$  do {

    while  $Tr_i < Gs_j$  do
         $Tr$  = next tuple in  $R$  after  $Tr$ ;                // continue scan of  $R$ 

    while  $Tr_i > Gs_j$  do
         $Gs$  = next tuple in  $S$  after  $Gs$                     // continue scan of  $S$ 

     $Ts = Gs$ ;                                              // Needed in case  $Tr_i \neq Gs_j$ 
    while  $Tr_i == Gs_j$  do {                               // process current  $R$  partition
         $Ts = Gs$ ;                                          // reset  $S$  partition scan
        while  $Ts_j == Tr_i$  do {                          // process current  $R$  tuple
            add  $\langle Tr, Ts \rangle$  to result;                // output joined tuples
             $Ts$  = next tuple in  $S$  after  $Ts$ ;              // advance  $S$  partition scan
        }  $Tr$  = next tuple in  $R$  after  $Tr$ ;                // advance scan of  $R$ 
    }                                                    // done with current  $R$  partition

     $Gs = Ts$ ;                                              // initialize search for next  $S$  partition
}
```




<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
36	lubber	6	36.0
44	guppy	5	35.0
58	rusty	10	35.0

Figure 12.9 An Instance of Sailors


<i>sid</i>	<i>bid</i>	<i>day</i>	<i>rname</i>
28	103	12/04/96	guppy
28	103	11/03/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Figure 12.10 An Instance of Reserves



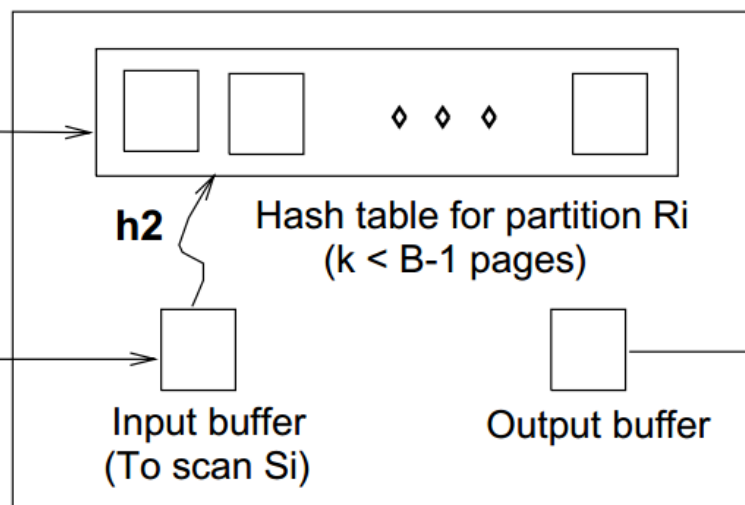
Probing Phase of Hash Join

Partitions of R and S




Disk

hash
function
h2



B main memory buffers

Join result



Disk

Figure 12.11 Probing Phase of Hash Join



```
// Partition  $R$  into  $k$  partitions
foreach tuple  $r \in R$  do
    read  $r$  and add it to buffer page  $h(r_i)$ ;           // flushed as page fills

// Partition  $S$  into  $k$  partitions
foreach tuple  $s \in S$  do
    read  $s$  and add it to buffer page  $h(s_j)$ ;           // flushed as page fills

// Probing Phase
for  $l = 1, \dots, k$  do {

    // Build in-memory hash table for  $R_l$ , using  $h_2$ 
    foreach tuple  $r \in$  partition  $R_l$  do
        read  $r$  and insert into hash table using  $h_2(r_i)$  ;

    // Scan  $S_l$  and probe for matching  $R_l$  tuples
    foreach tuple  $s \in$  partition  $S_l$  do {
        read  $s$  and probe table using  $h_2(s_j)$ ;
        for matching  $R$  tuples  $r$ , output  $\langle r, s \rangle$  };

    clear hash table to prepare for next partition;
}
```